



ECON408: Computational Methods in Macroeconomics

Stochastic Dynamics, AR(1) Processes, and Ergodicity

Jesse Perla

jesse.perla@ubc.ca

University of British Columbia

Table of contents

- Overview
- Random Variables Review
- Stochastic Processes
- AR(1) Processes
- Stationary Distributions
- Nonlinear Stochastic Processes
- Stochastic Growth Model
- Multiplicative Growth Processes

Overview

Motivation and Materials

- In this lecture, we will introduce our stochastic processes and review probability
- Our first example of a stochastic process is the **AR(1)** process (i.e. autoregressive of order one)
 - This is a simple, univariate process, but it is directly useful in many cases
- We will also introduce the concept of **ergodicity** to help us understand long-run behavior
- While this section is not directly introducing new economic models, it provides the backbone for our analysis of the wealth and income distribution

Deterministic Processes

- We have seen deterministic processes in previous lectures, e.g. the linear

$$X_{t+1} = aX_t + b$$

- These are coupled with an initial condition X_0 , which enables us to see the evolution of a variable
- The state variable, X_t , could be a vector
- The evolution could be non-linear $X_{t+1} = h(X_t)$, etc.
- But many states in the real world involve randomness

Materials

- Adapted from QuantEcon lectures coauthored with John Stachurski and Thomas J. Sargent
 - AR1 Processes
 - LLN and CLT
 - Continuous State Markov Chains

```
1 using LaTeXStrings, LinearAlgebra, Plots, Statistics
2 using Random, StatsPlots, Distributions, NLSolve
3 using Plots.PlotMeasures
4 default(;legendfontsize=16, linewidth=2, tickfontsize=12,
5           bottom_margin=15mm)
```

Random Variables Review

Random Variables


- **Random variables** are a collection of values with associated probabilities
- For example, a random variable Y could be the outcome of a coin flip
 - Let $Y = 1$ if heads and $Y = 0$ if tails
 - Assign probabilities $\mathbb{P}(Y = 1) = \mathbb{P}(Y = 0) = 0.5$
- or a **normal random variable** with mean μ and variance σ^2 , denoted $Y \sim \mathcal{N}(\mu, \sigma^2)$ has density $p(y) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y-\mu)^2}{2\sigma^2}\right)$

Discrete vs. Continuous Variables

- If discrete (e.g., $X \in \{x_1, \dots, x_N\}$), then
 - The **probability mass function** (pmf) is the probability of each value $p \in \mathbb{R}^N$
 - Such that $\sum_{i=1}^N p_i = 1$, and $p_i \geq 0$
 - i.e. $p_i = \mathbb{P}(X = x_i)$
- If continuous, then the **probability density function** (pdf) is the probability of each value and can be represented by a function
 - $p : \mathbb{R} \rightarrow \mathbb{R}$ if X is defined on \mathbb{R}
 - $\int_{-\infty}^{\infty} p(x)dx = 1$, and $p(x) \geq 0$
 - $\mathbb{P}(X = a) = 0$ in our examples, and $\mathbb{P}(X \in [a, b]) = \int_a^b p(x)dx$


Normal Random Variables

```
1 mu = 0.1
2 sigma = 0.5
3 d = Normal(0.1, sigma) # SD not variance
4 x = range(mu - 2 * sigma,
5           mu + 2 * sigma,
6           length=100)
7 plot(x, pdf.(d, x); label="Normal PDF",
8       xlabel="x", ylabel=L"p(x)",
9       size=(600,400))
```



Comparing to a Histogram

```
1 n = 1000
2 x_draws = rand(d, n) # gets n samples
3 histogram(x_draws; label="Histogram",
4           xlabel="x", ylabel=L"\hat{p}(x)",
5           normalize=true, size=(600,400))
6 plot!(x, pdf.(d, x); label="Normal PDF",
7        lw=3)
```



Normal Random Variables

- Normal random variables are special for many reasons (e.g., central limit theorems)
- They are the only continuous random variable with finite variance that is closed under linear combinations
 - For independent $X \sim \mathcal{N}(\mu_X, \sigma_X^2)$ and $Y \sim \mathcal{N}(\mu_Y, \sigma_Y^2)$
 - $aX + bY \sim \mathcal{N}(a\mu_X + b\mu_Y, a^2\sigma_X^2 + b^2\sigma_Y^2)$
 - Also true with multivariate normal distributions
- Common transformation taking out mean and variance
 - Could draw $Y \sim \mathcal{N}(\mu, \sigma^2)$
 - Or could draw $X \sim \mathcal{N}(0, 1)$ and then $Y = \mu + \sigma X$

Expectations

- For discrete-valued random variables

$$\mathbb{E}[f(X)] = \sum_{i=1}^N f(x_i)p_i$$

- For continuous valued random variables

$$\mathbb{E}[f(X)] = \int_{-\infty}^{\infty} f(x)p(x)dx$$

Moments

- The **mean** of a random variable is the first moment, $\mathbb{E}[X]$
- The **variance** of a random variable is the second moment, $\mathbb{E}[(X - \mathbb{E}[X])^2]$
 - Note the recentering by the mean. Could also calculate as
$$\mathbb{E}[X^2] - \mathbb{E}[X]^2$$
- Normal random variables are characterized by their first 2 moments

Law(s) of Large Numbers

- Let X_1, X_2, \dots be independent and identically distributed (iid) random variables with mean $\mu \equiv \mathbb{E}(X) < \infty$, then let

$$\bar{X}_n \equiv \frac{1}{n} \sum_{i=1}^n X_i$$

- One version is **Kolmogorov's Strong Law of Large Numbers**

$$\mathbb{P} \left(\lim_{n \rightarrow \infty} \bar{X}_n = \mu \right) = 1$$

→ i.e. the average of the random variables converges to the mean


Sampling and Plotting the Mean

```
1 function ksl(distribution, n = 100)
2     title = nameof(typeof(distribution))
3     observations = rand(distribution, n)
4     sample_means = cumsum(observations) ./ (1:n)
5     mu = mean(distribution)
6     plot(repeat((1:n)', 2), [zeros(1, n); observations']; title, xlabel="n",
7         label = "", color = :grey, alpha = 0.5)
8     plot!(1:n, observations; color = :grey, markershape = :circle,
9         alpha = 0.5, label = "", linewidth = 0)
10    if !isnan(mu)
11        hline!([mu], color = :black, linewidth = 1.5, linestyle = :dash,
12            grid = false, label = L"\mathbb{E}[X]")
13    end
14    return plot!(1:n, sample_means, linewidth = 3, alpha = 0.6, color = :green, label = L"\bar{X}_n")
15 end
```

ksl (generic function with 2 methods)

LLN with the Normal Distribution

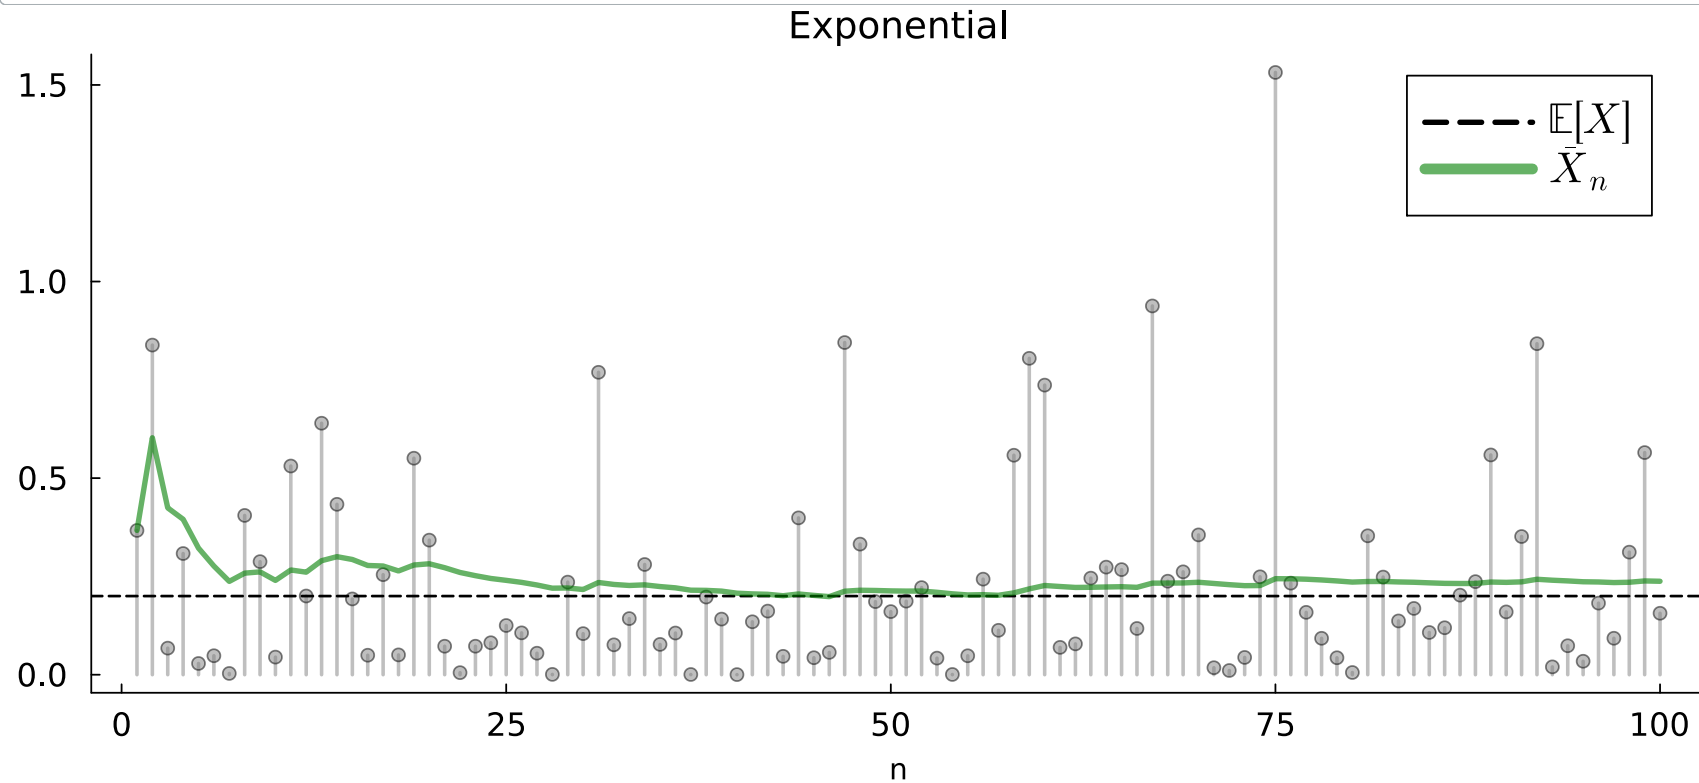
```
1 dist = Normal(0.0, 1.0) # unit normal
2 ksl(dist)
```



LLN with the Exponential

- $f(x) = \frac{1}{\alpha} \exp(-x/\alpha)$ for $x \geq 0$ with mean α


```
1 dist = Exponential(0.2)
2 ksl(dist)
```



LLN with the Cauchy?

- $f(x) = 1/(\pi(1 + x^2))$, with median = 0 and $\mathbb{E}(X)$ undefined

```
1 Random.seed!(0); # reproducible results
2 dist = Cauchy() # Doesn't have an expectation!
3 sample_mean = cumsum(rand(dist, n)) ./ (1:n)
4 plot(1:n, sample_mean, color = :red, alpha = 0.6, label = L"\bar{X}_n",
5      xlabel = "n", linewidth = 3)
6 hline!([0], color = :black, linestyle = :dash, label = "", grid = false)
```



Monte-Carlo Calculation of Expectations

- One application of this is the numerical calculation of expectations
- Let \mathbf{X} be a random variable with density $p(x)$, and hence $\mathbb{E}[f(\mathbf{X})] = \int_{-\infty}^{\infty} f(x)p(x)dx$ (or $\sum_{i=1}^N f(x_i)p_i$ if discrete)
- These integrals are often difficult to calculate analytically, but if we can draw $\mathbf{X} \sim p$, then we can approximate the expectation by

$$\mathbb{E}[f(\mathbf{X})] \approx \frac{1}{n} \sum_{i=1}^n f(x_i)$$

- Then by the LLN this converges to the true expectation as $n \rightarrow \infty$

Discrete Example

- Let X be a discrete random variable with N states and probabilities p_i
- Then $\mathbb{E}[f(X)] = \sum_{i=1}^N f(x_i)p_i$
- For example, the Binomial distribution and $f(x) = \log(x + 1)$

```
1 # number of trials and probability of success
2 dist = Binomial(10, 0.5)
3 plot(dist; label="Binomial PMF",
4       size=(600, 400))
5 vals = support(dist) # i.e. 0:10
6 p = pdf.(dist, vals)
7 # Calculate the expectation manually
8 @show mean(dist), dot(vals, p);
```

```
(mean(dist), dot(vals, p)) = (5.0,
5.0000000000000008)
```

Using Monte-Carlo

```

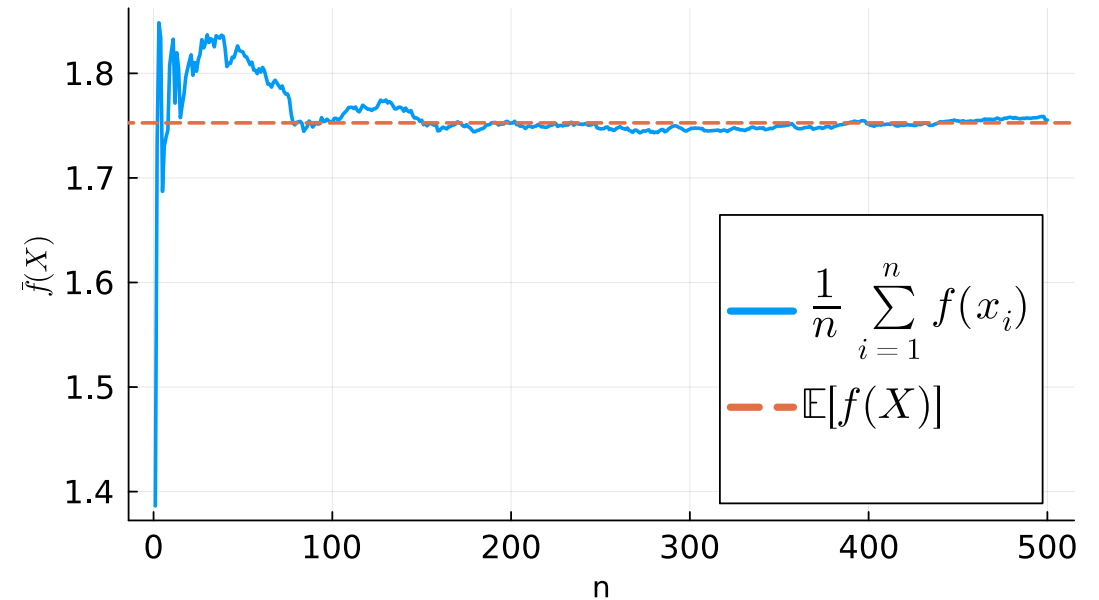
1 N = 500
2 # expectation with PMF, then MC
3 f_expect = dot(log.(vals .+ 1), p)
4 x_draws = rand(dist, N)
5 f_x_draws = log.(x_draws .+ 1)
6 f_expect_mc = sum(f_x_draws) / N
7 @show f_expect, f_expect_mc
8 # Just calculate sums then divide by N
9 f_means = cumsum(f_x_draws)./(1:N)
10 plot(1:length(f_means), f_means;
11      label=L"\frac{1}{n}\sum_{i=1}^n f(x_i)",
12      xlabel="n", ylabel=L"\bar{f}(X)",
13      size=(600,400))
14 hline!([f_expect];linestyle = :dash,
15        label = L"\mathbb{E}[f(X)]")

```

```

(f_expect, f_expect_mc) = (1.75263932077417,
1.7552834928857297)

```



Stochastic Processes

Stochastic Processes

- A **stochastic process** is a sequence of random variables
 - We will focus on **discrete time** stochastic processes, where the sequence is indexed by $t = 0, 1, 2, \dots$
 - Could be discrete or continuous random variables
- Skipping through some formality, assume that they share the same values but probabilities may change
- Denote then as a sequence $\{X_t\}_{t=0}^{\infty}$

Joint, Marginal, and Conditional Distributions

- Can ask questions on the probability distributions of the process
- The **joint distribution** of $\{X_t\}_{t=0}^{\infty}$ or a subset
 - In many cases things will be correlated over time or else no need to be a process
- The **marginal distribution** of X_t for any t
 - This is a proper PDF, marginalized from the joint distribution of all values
- **Conditional distributions**, fixing some values
 - e.g. X_{t+1} given X_t, X_{t-1} , etc. are known

Markov Process

- Before we go further, let's discuss a broader class of these processes useful in economics
- A **Markov process** is a stochastic process where the conditional distribution of \mathbf{X}_{t+1} given $\mathbf{X}_t, \mathbf{X}_{t-1}, \dots$ is the same as the conditional distribution of \mathbf{X}_{t+1} given \mathbf{X}_t
 - i.e. the future is independent of the past given the present
- Note that with the AR(1) model, if I know \mathbf{X}_t then I can calculate the PDF of \mathbf{X}_{t+1} directly without knowing the past
- This is “first-order” since only one lag is required, but could be higher order
 - A finite number of lags can always be added to the state vector to make it first-order

AR(1) Processes

A Simple Auto-Regressive Process with One Lag

$$X_{t+1} = aX_t + b + cW_{t+1}$$

- Just added randomness to the deterministic process from time t to $t + 1$
- $W_{t+1} \sim \mathcal{N}(0, 1)$ is IID “shocks” or “noise”
- Could have an initial condition for X_0 Or could have an initial distribution
 - X_t is a random variable, and so can X_0
 - “Degenerate random variable” if $P(X_0 = x) = 1$ for some x
 - Assume $X_0 \sim \mathcal{N}(\mu_0, v_0)$, and $v_0 \rightarrow 0$ is the degenerate case

Evolution of the AR(1) Process

- Both W_{t+1} and X_0 are assumed to be normally distributed
- As we discussed, linear combinations of normal random variables are normal
 - So X_t is normal for all t by induction
- Furthermore, we have a formula for the recursion
 - If $X_t \sim \mathcal{N}(\mu_t, v_t)$, then $X_{t+1} \sim \mathcal{N}(a\mu_t + b, a^2v_t + c^2)$
 - Hence, the evolution of the mean and variance follow a simple difference equation $\mu_{t+1} = a\mu_t + b$ and $v_{t+1} = a^2v_t + c^2$
 - Let $X_t \sim \psi_t \equiv \mathcal{N}(\mu_t, v_t)$


Visualizing the AR(1) Process

```
1 a = 0.9
2 b = 0.1
3 c = 0.5
4
5 # initial conditions mu_0, v_0
6 mu = -3.0
7 v = 0.6
```

0.6

Visualizing the AR(1) Process


```
1  sim_length = 5
2  x_grid = range(-5, 7, length = 120)
3
4  plt = plot(;size = (600, 400))
5  for t in 1:sim_length
6      mu = a * mu + b
7      v = a^2 * v + c^2
8      dist = Normal(mu, sqrt(v))
9      plot!(plt, x_grid, pdf.(dist, x_grid),
10         label = L"\psi_{%t}", linealpha = 0.7)
11  end
12  plt
```



From a Degenerate Initial condition

- Cannot plot ψ_0 since it is a point mass at μ_0

```
1 mu = -3.0
2 v = 0.0
3 plt = plot(;size = (600, 400))
4 for t in 1:sim_length
5     mu = a * mu + b
6     v = a^2 * v + c^2
7     dist = Normal(mu, sqrt(v))
8     plot!(plt, x_grid, pdf.(dist, x_grid),
9         label = L"\psi_{%t}", linealpha = 0.7)
10 end
11 plt
```



Practice with Iteration

- Let us practice creating a map and iterating it
- We will need to modify our `iterate_map` function to work with vectors
- Let $x \equiv [\mu \quad v]^\top$,

```
1 function iterate_map(f, x0, T)
2     x = zeros(length(x0), T + 1)
3     x[:, 1] = x0
4     for t in 2:(T + 1)
5         x[:, t] = f(x[:, t - 1])
6     end
7     return x
8 end
```

`iterate_map` (generic function with 1 method)

Implementation of the Recurrence for the AR(1)

```
1 function f(x;a, b, c)
2   mu = x[1]
3   v = x[2]
4   return [a * mu + b, a^2 * v + c^2]
5 end
6 x_0 = [-3.0, 0.6]
7 T = 5
8 x = iterate_map(x -> f(x; a, b, c), x_0, T)
```

2×6 Matrix{Float64}:

-3.0	-2.6	-2.24	-1.916	-1.6244	-1.36196
0.6	0.736	0.84616	0.93539	1.00767	1.06621

Using Matrices

$$x_{t+1} = \underbrace{\begin{bmatrix} a & 0 \\ 0 & a^2 \end{bmatrix}}_{\equiv A} x_t + \underbrace{\begin{bmatrix} b \\ c^2 \end{bmatrix}}_{\equiv B}$$

```
1 A = [a 0; 0 a^2]
2 B = [b; c^2]
3 x = iterate_map(x -> A * x + B, x_0, T)
```

2x6 Matrix{Float64}:

-3.0	-2.6	-2.24	-1.916	-1.6244	-1.36196
0.6	0.736	0.84616	0.93539	1.00767	1.06621

Fixed Point?

- Whenever you have maps, you can ask whether a fixed point exists
- This is especially easy to check here. Solve,
 - $\mu = a\mu + b \implies \mu = \frac{b}{1-a}$
 - $v = a^2v + c^2 \implies v = \frac{c^2}{1-a^2}$
- Lets check for a fixed point numerically

```
1 sol = fixedpoint(x -> A * x + B, x_0)
2 @show sol.zero
3 @show b/(1-a), c^2/(1-a^2);
```

```
sol.zero = [1.00000000000000249, 1.3157894736842035]
(b / (1 - a), c ^ 2 / (1 - a ^ 2)) = (1.0000000000000002, 1.3157894736842108)
```

Existence of a Fixed Point


- The importance of a is also clear when we look at the A matrix
- We know the eigenvalues of a diagonal matrix are the diagonal elements
→ i.e., $\lambda_1 = a$ and $\lambda_2 = a^2$
- If $|a| < 1$, then $a^2 < |a| < 1$ and hence the maximum absolute value of the eigenvalues is below 1
- As we saw in the univariate case, conditions of this sort were crucial to determine whether the systems would converge
- We will see more complicated versions of the A matrix as we move into richer “state space models”

Evolution of the Probability Distributions

```

1  x_0 = [-1.0, 0.1] # tight
2  T = 10
3  f(x) = A * x + B
4  x = iterate_map(f, x_0, T)
5  x_star = fixedpoint(f, x_0).zero
6  plt = plot(Normal(x_star[1], sqrt(x_star[2])));
7          label = L"\psi^*",
8          style = :dash,
9          size = (600, 400))
10 for t in 1:T
11     dist = Normal(x[1, t], sqrt(x[2, t]))
12     plot!(plt, dist, label = L"\psi_{%$(t-1)}",
13           linealpha = 0.7)
14 end
15 plt

```



Stationary Distributions

Fixed Points and Steady States

- Recall in the lecture on deterministic dynamics that we discussed fixed point and steady states $x_{t+1} = f(x_t)$ has a **fixed point** x^* if $x^* = f(x^*)$
 - e.g. $x_{t+1} = ax_t + b$ has $x^* = \frac{b}{1-a}$ if $|a| < 1$
- We can also interpret as a **steady state** x^* as $\lim_{t \rightarrow \infty} x_t = x^*$ for some x_0
 - Stability looked at stability which told us about which x^* the process would approach from points x_0 near x^*
- The key: for x^* if we apply $f(x^*)$ evolution equation and remain at that point

Stationary Distributions

- Analogously, with stochastic processes we can think about applying the evolution equation to random variables
 - Instead of a point, we have a distribution ψ^*
 - Then rather than checking $x^* = f(x^*)$, we check $\psi^* \sim f(\psi^*)$, where that notation is loosely taking into account the distribution of shocks
- Similar to stability, we can consider if repeatedly applying $f(\cdot)$ repeatedly to various ψ_0 converges to ψ^*

AR(1) Example

- Take $X_{t+1} = aX_t + b + cW_{t+1}$ if $|a| < 1$ for $W_{t+1} \sim \mathcal{N}(0, 1)$
- Recall If $X_t \sim \mathcal{N}(\mu_t, v_t) \equiv \psi_t$, then using properties of Normals
 - $X_{t+1} \sim \mathcal{N}(a\mu_t + b, a^2v_t + c^2) \equiv \psi_{t+1}$
 - We derived the fixed point of the mean and variance iteration as $\psi^* \sim \mathcal{N}(\mu^*, v^*) = \mathcal{N}\left(\frac{b}{1-a}, \frac{c^2}{1-a^2}\right)$
- Apply the evolution equation to ψ^* we demonstrate that $\psi^* \sim f(\psi^*)$

$$\mathcal{N}\left(a\frac{b}{1-a} + b, a^2\frac{c^2}{1-a^2} + c^2\right) = \mathcal{N}\left(\frac{b}{1-a}, \frac{c^2}{1-a^2}\right)$$


- i.e., from any initial condition, the distribution of X_t converges to ψ^*

What if $a > 1$?

```

1 a,b,c = 1.1, 0.2, 0.25
2 A = [a 0; 0 a^2]
3 B = [b; c^2]
4 f(x) = A * x + B
5 T = 15
6 x = iterate_map(f, [0.0, 0.1], T)
7 plt = plot(Normal(x[1, end], sqrt(x[2, end])));
8           label = L"\psi_{%$T}",
9           size = (600, 400))
10 for t in 1:5
11     dist = Normal(x[1, t], sqrt(x[2, t]))
12     plot!(plt, dist, label=L"\psi_{%$(t-1)}",
13           linealpha = 0.7)
14 end
15 plt

```



Analyzing the Failure of Convergence


- If it exists, the stationary distribution would need to be $\psi^* \equiv \mathcal{N} \left(\frac{b}{1-a}, \frac{c^2}{1-a^2} \right)$
- Note that if $b > 0$ we get the drift of the process forward
 - But, just as in the case of the deterministic process, this just acts as a force to move the distribution, not spread it out
- In fact, with $b = 0$ the mean of ψ_t is always 0, but the variance grows without bound if $c > 0$
- Lets plot the $a = 1, b = 0$ case

What if $a = 1, b = 0$?

```

1 a,b,c = 1.0, 0.0, 0.25
2 A = [a 0; 0 a^2]
3 B = [b; c^2]
4 f(x) = A * x + B
5 T = 15
6 x = iterate_map(f, [0.0, 0.1], T)
7 plt = plot(Normal(x[1, end], sqrt(x[2, end])));
8           label = L"\psi_{%$T}",
9           size = (600, 400))
10 for t in 1:5
11     dist = Normal(x[1, t], sqrt(x[2, t]))
12     plot!(plt, dist, label=L"\psi_{%$(t-1)}",
13           linealpha = 0.7)
14 end
15 plt

```



Ergodicity

- There are many different variations and definitions of ergodicity
- Among other things, this rules out are cases where the process is “trapped” in a subset of the state space and can’t switch out
- Also ensures that the distribution doesn’t spread or drift asymptotically
- Ergodicity lets us apply LLNs to the stochastic process, even though they are not independent

Ergodicity

- We will consider a process $\{X_t\}_{t=0}^{\infty}$ with a stationary distribution ψ^*
- The process is **ergodic** if for any $f : \mathbb{R} \rightarrow \mathbb{R}$ (with regularity conditions)

$$\lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=1}^T f(X_t) = \int f(x) \psi^*(x) dx$$

- i.e. the time average of the function converges to the expectation of the function. Mean ergodic if only require this to work for $f(x) = x$

Iteration with IID Noise

- Adapt scalar iteration for iid noise

```
1 function iterate_map_iid(f, dist, x0, T)
2     x = zeros(T + 1)
3     x[1] = x0
4     for t in 2:(T + 1)
5         x[t] = f(x[t - 1], rand(dist))
6     end
7     return x
8 end
9 a,b,c = 0.9, 0.1, 0.05
10 x_0 = 0.5
11 T = 5
12 h(x, W) = a * x + b + c * W # iterate given random shock
13 x = iterate_map_iid(h, Normal(), x_0, T)
```

6-element Vector{Float64}:


```
0.5
0.5252717486805177
0.5306225876900339
0.46819901566492783
0.532032538532688
0.583020976850554
```


Demonstration of Ergodicity with Mean

```

1 T = 2000
2 x_0 = 0.5
3 x = iterate_map_iid(h, Normal(), x_0, T)
4 x_means = cumsum(x)./(1:(T+1))
5 plot(0:T, x_means;
6     label=L"\frac{1}{t}\sum_{s=0}^{t-1} X_s",
7     xlabel = "t", size = (600, 400))
8 hline!([b/(1-a)], color = :black,
9     linestyle = :dash,
10    label = L"\mathbb{E}[X]")

```




Starting at the Stationary Distribution

- A reasonable place to start many simulations is a draw from the stationary distribution

```

1 Random.seed!(20)
2 x_0 = rand(Normal(b/(1-a), sqrt(c^2/(1-a^2))))
3 x = iterate_map_iid(h, Normal(), x_0, T)
4 x_means = cumsum(x)./(1:(T+1))
5 plot(0:T, x_means;
6     label=L"\frac{1}{t}\sum_{s=0}^{t-1} X_s",
7     xlabel = "t", size = (600, 400))
8 hline!([b/(1-a)], color = :black,
9     linestyle = :dash,
10    label = L"\mathbb{E}[X]")

```



The Speed of Convergence


- The speed with which the process converges towards its stationary distribution is important
- Key things which govern this transition will be
 - Autocorrelation: As a goes closer to 0 , the faster it converges back towards the mean - as with deterministic processes
 - Variances: With large c the noise may dominate and the ψ^* becomes broader

Close to a Random Walk

```

1 Random.seed!(20)
2 a,b,c = 0.99, 0.01, 0.05
3 h(x, W) = a * x + b + c * W
4 T = 2000
5 x_0 = 0.5
6 x = iterate_map_iid(h, Normal(), x_0, T)
7 x_means = cumsum(x)./(1:(T+1))
8 plot(0:T, x_means;
9     label=L"\frac{1}{t}\sum_{s=0}^{t-1} X_s",
10    xlabel = "t", size = (600, 400))
11 hline!([b/(1-a)], color = :black,
12        linestyle = :dash,
13        label = L"\mathbb{E}[X]")

```



Dependence on Initial Condition

- Intuition: ergodicity is that the initial conditions “wear off” over time
- However, even if a process is ergodic and has a well-defined stationary distribution, it may take a long time to converge to it
- This is very important in many quantitative models:
 - How much does your initial wealth matter for your long-run?
 - If your wages start low due to discrimination, migration, or just bad luck, how long does it converge?
 - If we provide subsidies to new firms, how long would it take for that to affect the distribution of firms?

Example of a Non-Ergodic Stochastic Process

- Between $t = 0$ and $t = 1$ a coin is flipped (e.g., result of key exam)
 - If heads: income follows $X_{t+1} = aX_t + b + cW_{t+1}$ with $b = 0.1$ for $t \geq 1$
 - If tails: income follows $X_{t+1} = aX_t + b + cW_{t+1}$ with $b = 1.0$ for $t \geq 1$
- The initial condition and early sequence cannot be forgotten
- If there is ANY probability of switching between careers, then it is ergodic because it “mixes”

Moving Average Representation, $MA(\infty)$, for $AR(1)$

- From $X_t = aX_{t-1} + b + cW_t$, iterate backwards to X_0 and W_1

$$\begin{aligned} X_t &= a(aX_{t-2} + b + cW_{t-1}) + b + cW_t \\ &= a^2X_{t-2} + b(1 + a) + c(W_t + aW_{t-1}) \\ &= a^2(aX_{t-3} + b + cW_{t-2}) + b(1 + a) + c(W_t + aW_{t-1}) \\ &= a^tX_0 + b \sum_{j=0}^{t-1} a^j + c \sum_{j=0}^{t-1} a^j W_{t-j} \\ &= a^tX_0 + b \frac{1 - a^t}{1 - a} + c \sum_{j=0}^{t-1} a^j W_{t-j} \end{aligned}$$

Interpreting the Auto-Regressive Parameter

- The distribution of \mathbf{X}_t then depends on the distribution of \mathbf{X}_0 and the distribution of the sum of $t - 1$ iid random variables
- If \mathbf{X}_0 and \mathbf{W}_t are normal, then \mathbf{X}_t is normal since it is a linear combination

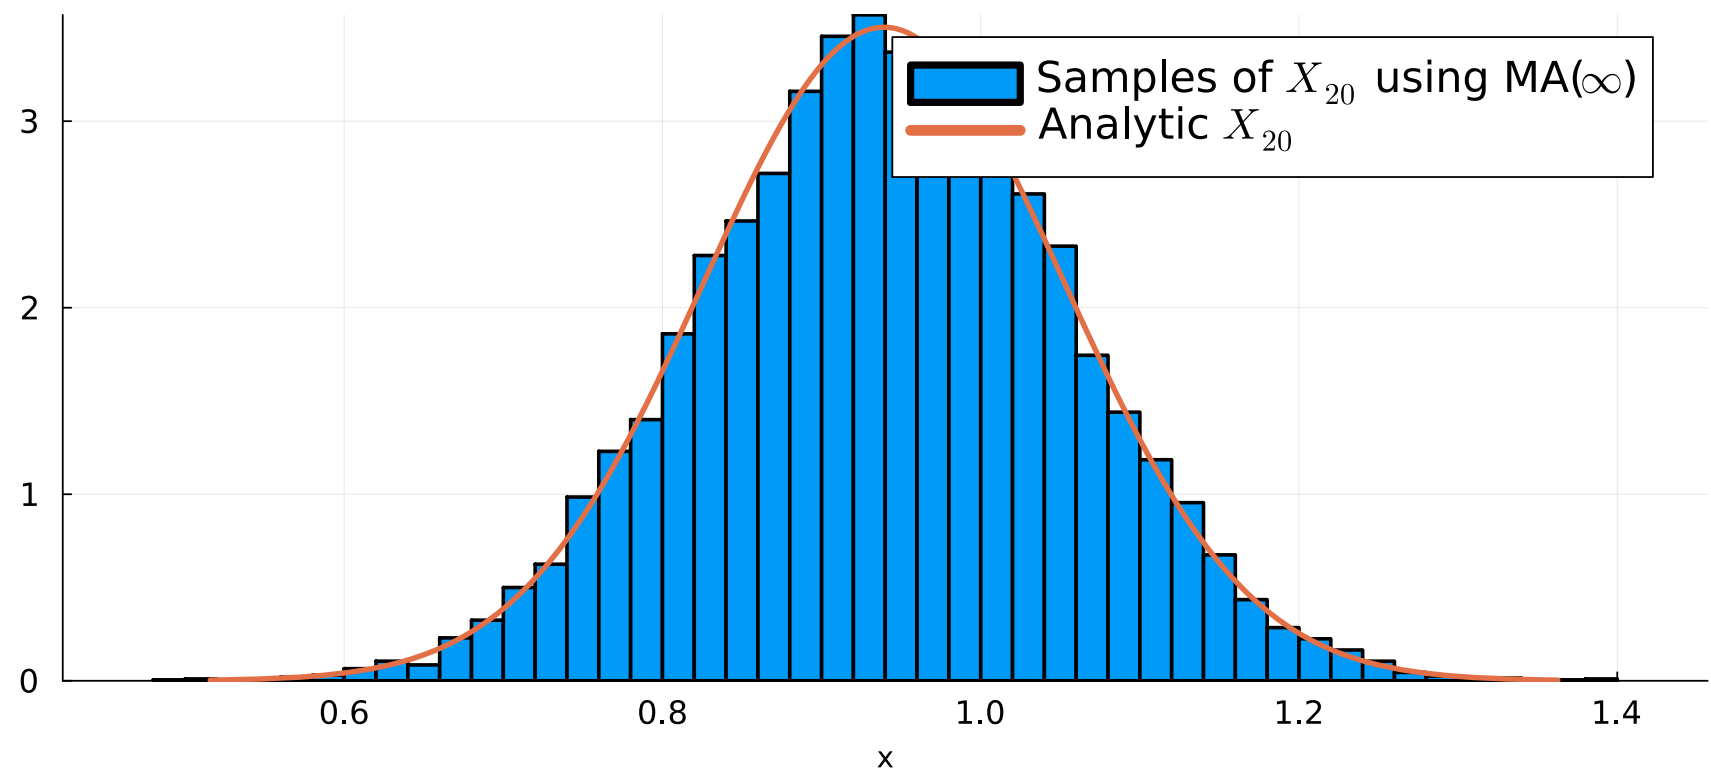
$$\mathbf{X}_t = a^t \mathbf{X}_0 + b \frac{1 - a^t}{1 - a} + c \sum_{j=0}^{t-1} a^j \mathbf{W}_{t-j}$$

- If $a = 1$ then the initial condition is never “forgotten”
- If $a = 1$, \mathbf{W}_{t-j} shocks are just as important determining the distribution of \mathbf{X}_t because the a^2 doesn’t “decay” over time

Simulation of Moving Average Representation

```
1 X_0 = 0.5 # degenerate prior
2 a, b, c = 0.9, 0.1, 0.05
3 A = [a 0; 0 a^2]
4 B = [b; c^2]
5 T = 20
6 num_samples = 10000
7 Xs = iterate_map(x -> A * x + B, [X_0, 0], T)
8 X_T = Normal(Xs[1, end], sqrt(Xs[2, end]))
9 W = randn(num_samples, T)
10 # Comprehensions and generators example, looks like math
11 X_T_samples = [a^T * X_0 + b * (1-a^T)/(1-a) + c * sum(a^j * W[i, T-j] for j in 0:T-1)
12                 for i in 1:num_samples]
13 histogram(X_T_samples; xlabel="x", normalize=true,
14           label=L"Samples of $X_{%T}$ using MA($\infty$)")
15 plot!(X_T; label=L"Analytic $X_{%T}$", lw=3)
```

Simulation of Moving Average Representation





Nonlinear Stochastic Processes

Nonlinearity with Additive Shocks

- A useful class involves nonlinear functions for the drift and variance

$$X_{t+1} = \mu(X_t) + \sigma(X_t)W_{t+1}$$

- IID W_{t+1} with $\mathbb{E}[W_{t+1}] = 0$ and frequently $\mathbb{E}[W_{t+1}^2] = 1$
- Nests our AR(1) process
 - $\mu(x) = ax + b$ and $\sigma(x) = c$

Auto-Regressive Conditional Heteroskedasticity (ARCH)

- For example, we may find that time-series data has time-varying volatility and depends on 1 lags


$$X_{t+1} = aX_t + \sigma_t W_{t+1}$$

- And that the variance increases as we move away from the mean of the stationary distribution $\sigma_t^2 = \beta + \gamma X_t^2$
- Hence the process becomes an ARCH(1)

$$X_{t+1} = aX_t + (\beta + \gamma X_t^2)^{1/2} W_{t+1}$$

Simulation of ARCH(1)

```
1 a = 0.7
2 beta, gamma = 5, 0.5
3 X_0 = 1.0
4 T = 200
5 h(x, W) = a * x + sqrt(beta + gamma * x^2) * W
6 x = iterate_map_iid(h, Normal(), X_0, T)
7 plot(x; label = L"X_t", size = (600, 400))
```



AR(1) with a Barrier


- Nonlinearity in economics often comes in various forms of barriers, e.g. borrowing constraints
- Consider our AR(1) except that the process can never go below **0**

$$X_{t+1} = \max\{aX_t + b + cW_{t+1}, 0.0\}$$

- We could **stop** the process at this point, but instead we will continue to iterate

Simulation of AR(1) with a Barrier


```
1 a,b,c = 0.95, 0.00, 0.05
2 X_min = 0.0
3 h(x, W) = max(a * x + b + c * W, X_min)
4 T = 1000
5 x_0 = 0.5
6 x = iterate_map_iid(h, Normal(), x_0, T)
7 plot(x; label = L"X_t", size = (600, 400))
```



Histogram of the AR(1) with a Barrier

- There isn't a true density of ψ^* due to the point mass at 0

```
1 T = 20000
2 x = iterate_map_iid(h, Normal(), x_0, T)
3 histogram(x; label = L"X_t", normalize = true,
4           xlabel = "x", size = (600, 400))
```



Stochastic Growth Model

Simple Growth Model with Stochastic Productivity

- Turning off population growth, for $f(k) = k^\alpha$, and s, δ constants

$$k_{t+1} = (1 - \delta)k_t + sZ_t f(k_t), \quad \text{given } k_0$$

- Let log productivity, $z_t \equiv \log Z_t$, follow an AR(1) process (why logs?)

$$\log Z_{t+1} = a \log Z_t + b + cW_{t+1}$$


Stationary Distribution of Productivity

- Recall that the stationary distribution of $\log Z_t$ is $\mathcal{N}\left(\frac{b}{1-a}, \frac{c^2}{1-a^2}\right)$
- Given the stationary distribution of Z_t is lognormal, we can check ergodicity

```

1 a, b, c = 0.9, 0.1, 0.05
2 Z_0 = 1.0
3 T = 20000
4 h(z, W) = a * z + b + c * W
5 z = iterate_map_iid(h, Normal(), log(Z_0), T)
6 Z = exp.(z)
7 histogram(Z; label = L"Z_t", normalize = true,
8           xlabel = "Z", size = (600, 400))
9 plot!(LogNormal(b/(1-a), sqrt(c^2/(1-a^2))),
10       lw = 3, label = L"\psi^*")

```



Quantiles

- Reminder: A quantile q is the x such that $\mathbb{P}(X \leq x) = q$
- Or, given a density $f(x)$ the quantile is the x such that $\int_{-\infty}^x f(x)dx = q$
- With data we can calculate an empirical quantile by first sorting the data, then finding the value of the observations below a certain count which is the proportion of the elements
 - e.g. with 100 observations, the 5th percentile is the 5th smallest observation
- The 0.5 quantile (i.e., the 50th percentile) is the median
- For heavily skewed distributions, the median is often a better measure of central tendency than the mean

Practice with Iteration and Multivariate Functions


```
1 function iterate_map_iid_vec(h, dist, x0, T)
2     x = zeros(length(x0), T + 1)
3     x[:, 1] = x0
4     for t in 2:(T + 1)
5         # accepts whatever type rand(dist) returns
6         x[:, t] = h(x[:, t - 1], rand(dist))
7     end
8     return x
9 end
```

iterate_map_iid_vec (generic function with 1 method)

Simulation of the Stochastic Growth Model

```
1 alpha, delta, s = 0.3, 0.1, 0.2
2 a, b, c = 0.9, 0.1, 0.1
3 function h(x, W)
4     k = x[1]
5     z = x[2]
6     return [(1-delta) * k + s * exp(z) * k^alpha,
7             a * z + b + c * W]
8 end
9 x_0 = [7.0, log(2.0)] # k_0, z_0
10 T = 150
11 x = iterate_map_iid_vec(h, Normal(), x_0, T)
12 plot(x[1, :]; label = L"k_t", xlabel = "t", size = (600, 400), legend=:topright)
13 plot!(exp.(x[2, :]), label = L"Z_t")
14 dist = LogNormal(b/(1-a), sqrt(c^2/(1-a^2)))
15 hline!([mean(dist)]; linestyle = :dash, label = L"\mathbb{E}[Z_t]")
16 hline!([quantile(dist, 0.05)]; lw=0, fillrange = [quantile(dist, 0.95)], fillalpha=0.2, label = "5/95")
```


Simulation of the Stochastic Growth Model



Ergodicity and Capital Accumulation

- Evaluate the closed-form steady-state capital k^* for the deterministic model

```
1 # Remember nonstochastic steady-state
2 k_ss_det= (s*mean(dist)/delta)^(1/(1-alpha))
3
4 T = 200000
5 x = iterate_map_iid_vec(h, Normal(), x_0, T)
6 histogram(x[1, :]; label = L"k_t",
7           normalize = true, xlabel = "k",
8           alpha=0.5, size = (600, 400))
9 vline!([k_ss_det]; linestyle = :dash, lw=3,
10        label = L"k_{ss}(c = 0)")
```



Multiplicative Growth Processes

Proportional Growth

- Many values grow or shrink proportional to their current size
 - e.g. population, firms, wealth, etc.
- The growth rates are themselves often random
 - e.g. population growth rates, firm growth rates, returns on wealth
 - Random good or bad luck can compound, which changes the distribution
- See [here](#) for more

Kesten Process

- The simplest **Kesten Process** is a process of the form

$$X_{t+1} = a_{t+1}X_t + y_{t+1}$$

- X_t is a state variable
- a_{t+1} is an IID random growth rate
- y_{t+1} is an IID random shock
- Examples: if population is N_t and growth rate between t and $t + 1$ is g_{t+1}
 - Then $N_{t+1}/N_t = 1 + g_{t+1}$
 - If we had migration y_{t+1} , then $N_{t+1} = (1 + g_{t+1})N_t + y_{t+1}$
- Key questions will be about whether stationary distributions exist, how they depend on parameters, and how fast they are approached

Conditions for a Stationary Distribution

- A stationary distribution may not exist.
- Important conditions for stationary are that
 - $\mathbb{E}[\log a_t] < 0$, intuition: $a_t < 1$ most of the time
 - $\mathbb{E}[y_t] < \infty$
- See [Kesten Processes](#) for more

Example with Random Growth on a Asset

- Let R_t be the gross returns on a asset, and W_t be value of it

$$W_{t+1} = R_{t+1}W_t$$

- i.e. no additional savings or consumption
- Let $\log R_t \sim \mathcal{N}(\mu, \sigma^2)$, i.e. lognormally distributed
 - The support of R_t is $(0, \infty)$ and $\mathbb{E}(R_t) = \exp(\mu + \sigma^2/2)$

Simulation

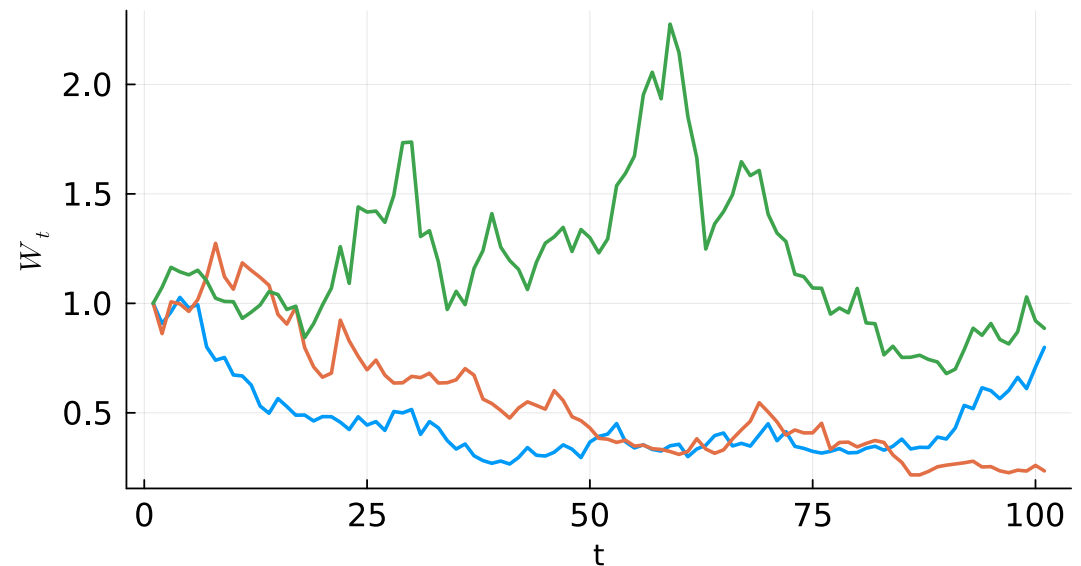
```

1 mu = -0.01
2 sigma = 0.1
3 R_dist = LogNormal(mu, sigma)
4 T = 100
5 W_0 = 1.0
6 @show mean(R_dist)
7 @show exp(mu + sigma^2/2)
8 plot(iterate_map_iid((W, R) -> W * R, R_dist,
9                      W_0, T);
10      ylabel = L"W_t", xlabel = "t",
11      size = (600, 400), legend=nothing,
12      title = "Simulations of Value")
13 plot!(iterate_map_iid((W, R) -> W * R, R_dist,
14                       W_0, T))
15 plot!(iterate_map_iid((W, R) -> W * R, R_dist,
16                       W_0, T))

```

$\text{mean}(R_dist) = 0.9950124791926823$
 $\exp(\mu + \sigma^2 / 2) = 0.9950124791926823$

Simulations of Value




Simulating an Ensemble

- Frequently we will want to simulate a large number of paths

```
1 function iterate_map_iid_ensemble(f, dist, x0, T, num_samples)
2     x = zeros(num_samples, T + 1)
3     x[:, 1] .= x0
4     for t in 2:(T + 1)
5         # or could do a loop over samples
6         x[:, t] .= f.(x[:, t - 1], rand(dist, num_samples))
7     end
8     return x
9 end
10 num_samples = 200
11 W = iterate_map_iid_ensemble((W, R) -> W * R, R_dist, W_0, T, num_samples)
12 plot(W'; ylabel = L"W_t", xlabel = "t", legend = nothing, alpha = 0.1,
13     color=:blue, lw = 1)
```

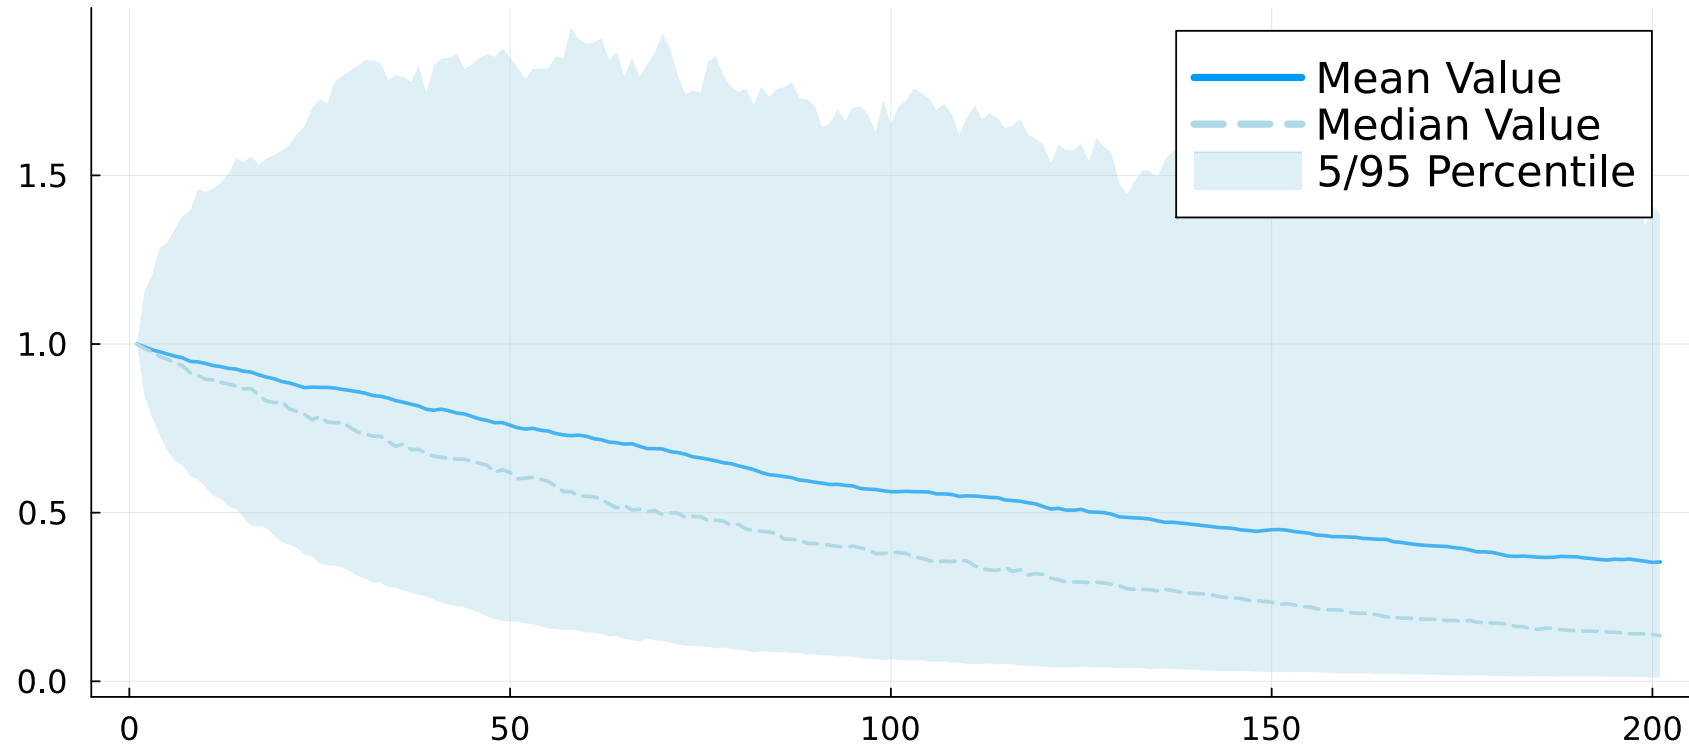

Simulating an Ensemble



Displaying the Distribution Dynamics

```
1 num_samples = 1000
2 T = 200
3 W = iterate_map_iid_ensemble((W, R) -> W * R, R_dist, W_0, T, num_samples)
4 q_50 = [quantile(W[:,i], 0.5) for i in 1:T+1]
5 q_05 = [quantile(W[:,i], 0.05) for i in 1:T+1]
6 q_95 = [quantile(W[:,i], 0.95) for i in 1:T+1]
7 mean_W = mean(W, dims=1)'
8 plot(mean_W; label = "Mean Value")
9 plot!(q_50; label = "Median Value", style = :dash, color = :lightblue)
10 plot!(q_05; label = "5/95 Percentile", lw=0, fillrange = q_95, fillalpha=0.4, color = :lightblue)
```

Displaying the Distribution Dynamics




Larger Returns

```
1 mu = -0.001
2 sigma = 0.1
3 R_dist = LogNormal(mu, sigma)
4 T = 100
5 W_0 = 1.0
6 @show mean(R_dist)
7 num_samples = 1000
8 W = iterate_map_iid_ensemble((W, R) -> W * R, R_dist, W_0, T, num_samples)
9 q_50 = [quantile(W[:,i], 0.5) for i in 1:T+1]
10 q_05 = [quantile(W[:,i], 0.05) for i in 1:T+1]
11 q_95 = [quantile(W[:,i], 0.95) for i in 1:T+1]
12 mean_W = mean(W, dims=1)'
13 plot(mean_W; label = "Mean Value")
14 plot!(q_50; label = "Median Value", style = :dash, color = :lightblue)
15 plot!(q_05; label = "5/95 Percentile", lw=0, fillrange = q_95, fillalpha=0.4, color = :lightblue)
```

Larger Returns

`mean(R_dist) = 1.004008010677342`



Divergence and Tails of Distributions

- These examples show that for multiplicative processes the distributions will often fan out, and potentially diverge
- This is a common feature of many economic and financial time series
- In particular, theory will show that for Kesten Processes, the tails of the distribution will be heavy even if it converges to a stationary distribution
 - i.e. the probability of large deviations from the mean will be higher than for a normal distribution
 - These will have what we call Power Law tails in the next section